

Enron Data Person of Interest Identification

Import packages and dataset

```
In [2]: # Initialize with imports
import sys
import pickle
sys.path.append("../tools/")

from feature_format import featureFormat, targetFeatureSplit
from tester import dump_classifier_and_data, main, load_classifier_and_data, test_classifier

import numpy as np
import pandas as pd
import future
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: ### Task 1: Select what features you'll use.
### features_list is a list of strings, each of which is a feature name.
### The first feature must be "poi".
features_list = ['poi', 'salary'] # You will need to use more features
```

```
In [4]: ### Load the dictionary containing the dataset
with open("final_project_dataset.pkl", "r") as data_file:
    data_dict = pickle.load(data_file)
```

Exploration


```
In [5]: # Remove outliers found below
del data_dict['TOTAL']
```

```
In [6]: # Make dict into df and print first 5 rows
df = pd.DataFrame.from_dict(data_dict, orient='index')
df.head()
```

```
Out[6]:
```

	salary	to_messages	deferral_payments	total_payments	exercised_stock_options
ALLEN PHILLIP K	201955	2902	2869717	4484442	1729541
BADUM JAMES P	NaN	NaN	178980	182466	257817
BANNANTINE JAMES M	477	566	NaN	916197	4046157
BAXTER JOHN C	267102	NaN	1295738	5634343	6680544
BAY FRANKLIN R	239671	NaN	260455	827696	NaN

5 rows × 21 columns



```
In [7]: # Print df info
df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 145 entries, ALLEN PHILLIP K to YEAP SOON
Data columns (total 21 columns):
salary                145 non-null object
to_messages           145 non-null object
deferral_payments     145 non-null object
total_payments        145 non-null object
exercised_stock_options 145 non-null object
bonus                 145 non-null object
restricted_stock       145 non-null object
shared_receipt_with_poi 145 non-null object
restricted_stock_deferred 145 non-null object
total_stock_value     145 non-null object
expenses              145 non-null object
loan_advances         145 non-null object
from_messages         145 non-null object
other                 145 non-null object
from_this_person_to_poi 145 non-null object
poi                   145 non-null bool
director_fees         145 non-null object
deferred_income       145 non-null object
long_term_incentive   145 non-null object
email_address         145 non-null object
from_poi_to_this_person 145 non-null object
dtypes: bool(1), object(20)
memory usage: 23.9+ KB
```

```
In [8]: # Make list of all names (keys) in dict and alphabetize
names = list(df.index.values)
```

```
In [9]: # Make list of all features (sub keys) in dict and print it
all_features = list(df.columns)
all_features
```

```
Out[9]: ['salary',
         'to_messages',
         'deferral_payments',
         'total_payments',
         'exercised_stock_options',
         'bonus',
         'restricted_stock',
         'shared_receipt_with_poi',
         'restricted_stock_deferred',
         'total_stock_value',
         'expenses',
         'loan_advances',
         'from_messages',
         'other',
         'from_this_person_to_poi',
         'poi',
         'director_fees',
         'deferred_income',
         'long_term_incentive',
         'email_address',
         'from_poi_to_this_person']
```

```
In [10]: # Replace 'NaN' values with actual nulls
df[all_features] = df[all_features].replace('NaN', np.nan)
```

```
In [11]: # Make list of POIs and print
pois = df[df['poi'] == True].index.tolist()
pois
```

```
Out[11]: ['BELDEN TIMOTHY N',
         'BOWEN JR RAYMOND M',
         'CALGER CHRISTOPHER F',
         'CAUSEY RICHARD A',
         'COLWELL WESLEY',
         'DELAINEY DAVID W',
         'FASTOW ANDREW S',
         'GLISAN JR BEN F',
         'HANNON KEVIN P',
         'HIRKO JOSEPH',
         'KOENIG MARK E',
         'KOPPER MICHAEL J',
         'LAY KENNETH L',
         'RICE KENNETH D',
         'RIEKER PAULA H',
         'SHELBY REX',
         'SKILLING JEFFREY K',
         'YEAGER F SCOTT']
```

```
In [12]: # Size of dataset
print("Number of people in dict: ", len(names))
print("Number of features in dict: ", len(all_features))
print("Number of POIs in dict: ", len(pois))
```

```
Number of people in dict: 145
Number of features in dict: 21
Number of POIs in dict: 18
```

```
In [13]: # Find number of NaNs for each feature
df.isnull().sum()
```

















```
Out[13]: salary                    51
to_messages                      59
deferral_payments              107
total_payments                  21
exercised_stock_options         44
bonus                           64
restricted_stock                 36
shared_receipt_with_poi         59
restricted_stock_deferred       128
total_stock_value                20
expenses                         51
loan_advances                   142
from_messages                    59
other                           53
from_this_person_to_poi         59
poi                              0
director_fees                   129
deferred_income                  97
long_term_incentive             80
email_address                   34
from_poi_to_this_person         59
dtype: int64
```







```
In [14]: # Make df with just pois and use to find number of NaNs for pois for each feature and print list
poi_df = df[df['poi'] == True]
poi_df.isnull().sum()
```

```
Out[14]: salary          1
to_messages            4
deferral_payments     13
total_payments         0
exercised_stock_options 6
bonus                  2
restricted_stock        1
shared_receipt_with_poi 4
restricted_stock_deferred 18
total_stock_value       0
expenses               0
loan_advances          17
from_messages           4
other                   0
from_this_person_to_poi 4
poi                     0
director_fees           18
deferred_income         7
long_term_incentive     6
email_address           0
from_poi_to_this_person 4
dtype: int64
```

```
In [15]: # Make df without pois
non_poi_df = df[df['poi'] == False]
```

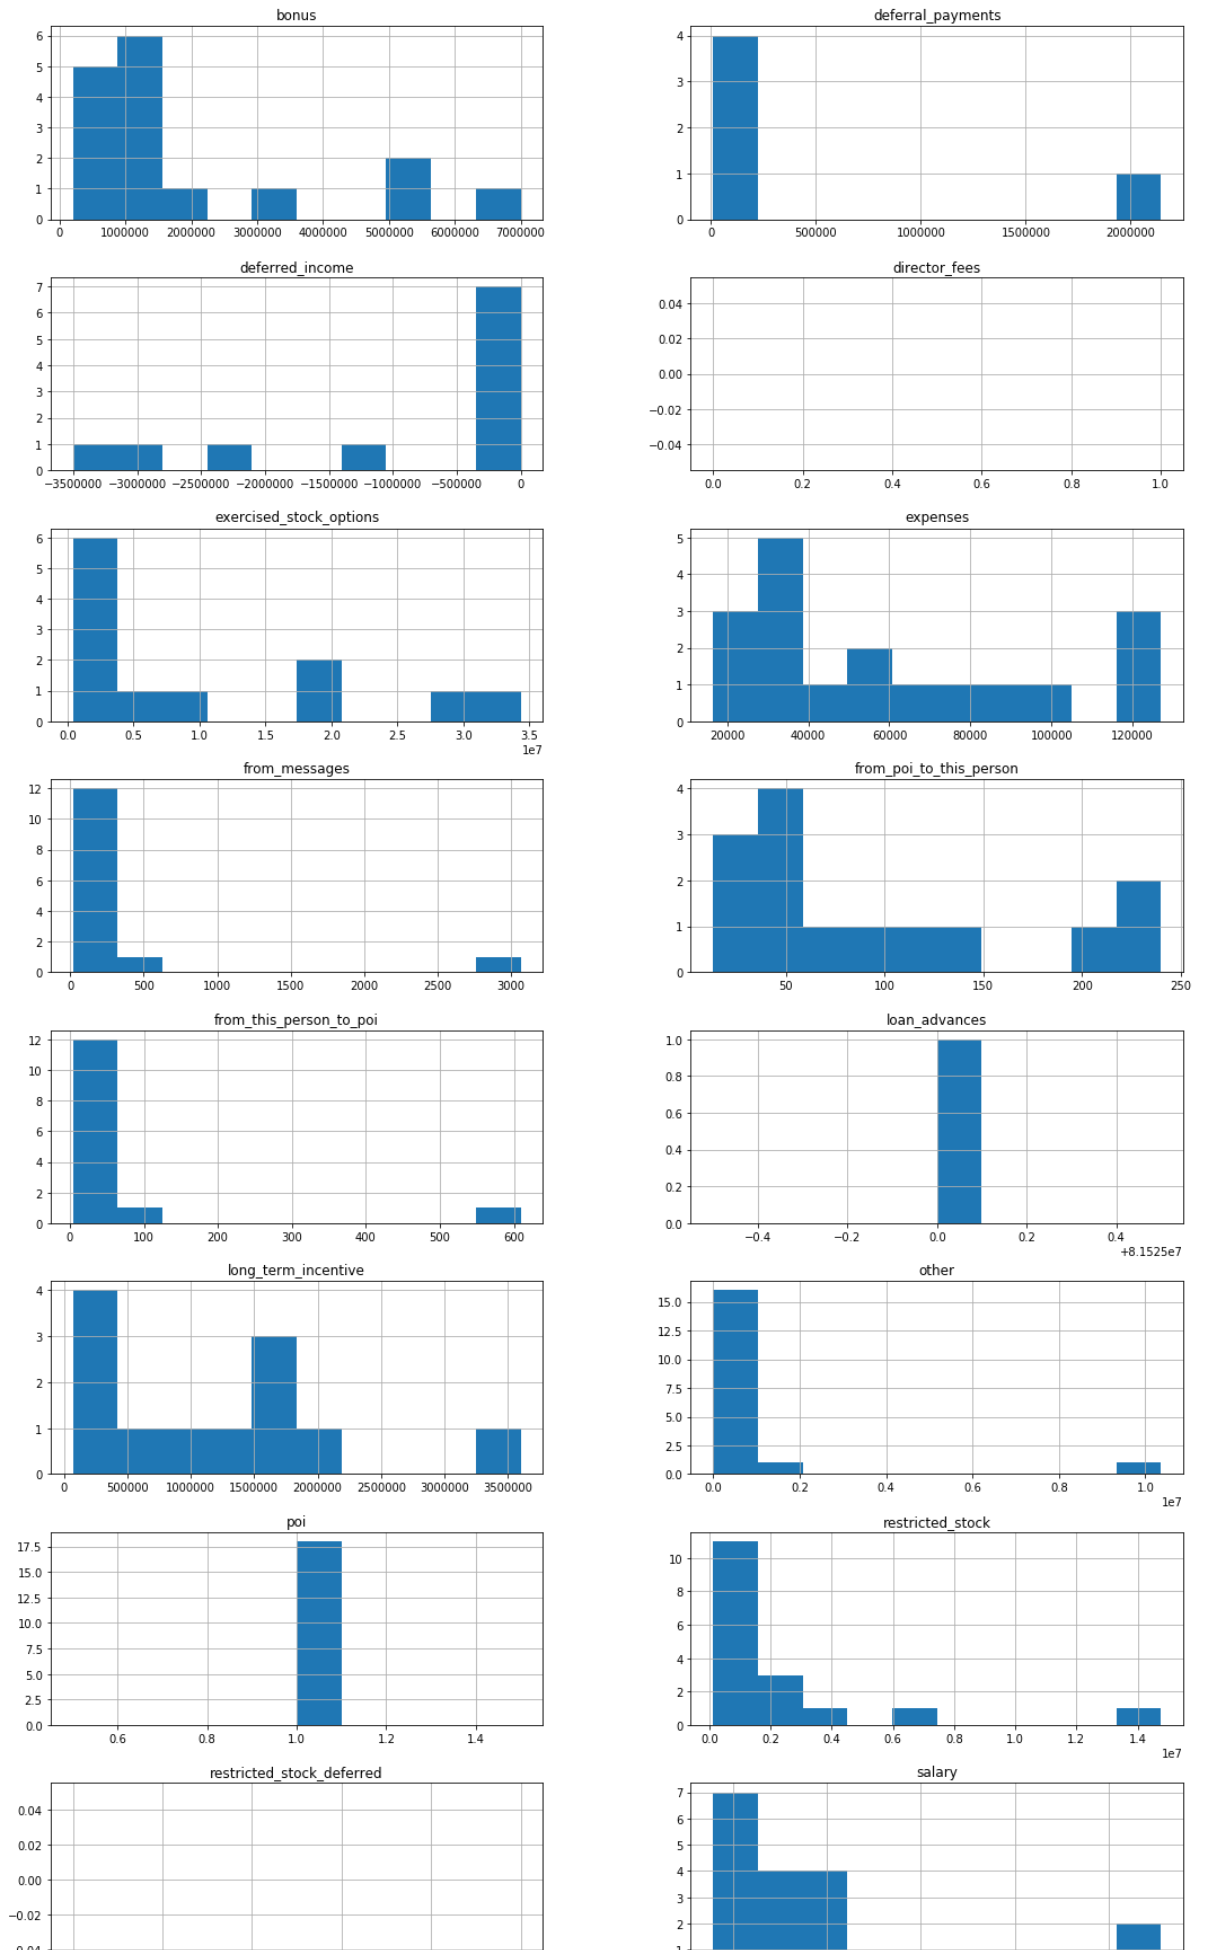
```
In [16]: # Plot histograms of each features distribution  
hist = df.hist(figsize=(18,40), layout=(10,2))
```

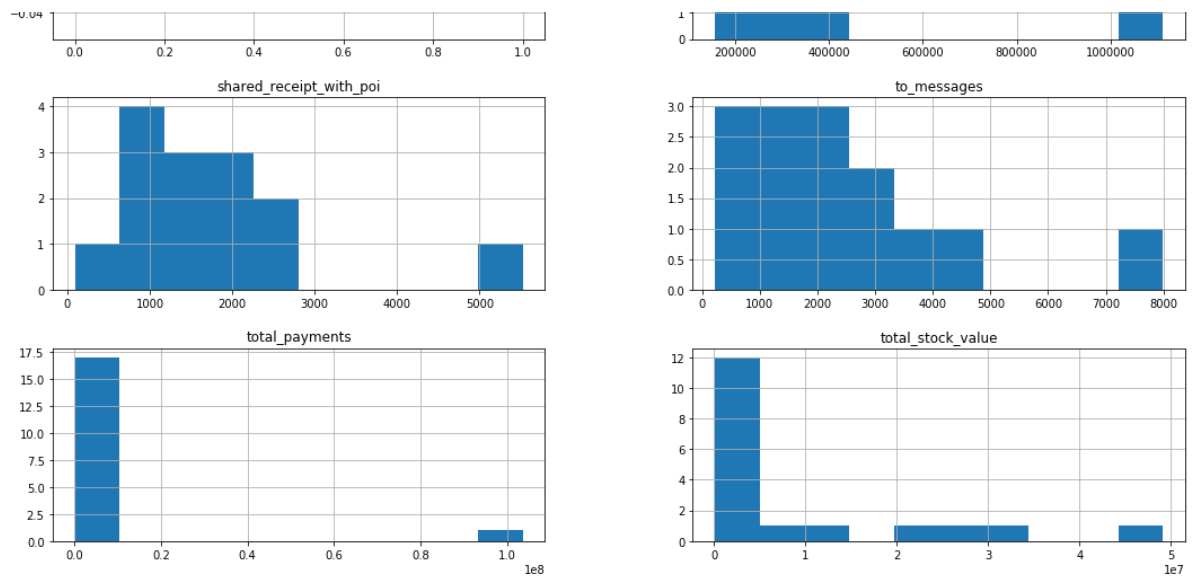





Data_dict contains key 'TOTAL' which is the total value for each feature and should not be included. Go back and remove this key from dict.


```
In [17]: # Plot histograms of each features distribution for pois  
hist = poi_df.hist(figsize=(18,40), layout=(10,2))
```







```
In [18]: # Non POI
# Plot histograms of each features distribution
hist = non_poi_df.hist(column=['salary', 'to_messages', 'deferral_payments',
                              'total_payments'],
                        figsize=(18,3), layout=(1,4))
```




```
In [19]: # POI
# Plot histograms of each features distribution
hist = poi_df.hist(column=['salary', 'to_messages', 'deferral_payments', 'total_payments'],
                    figsize=(18,3), layout=(1,4))
```




```
In [20]: # Non POI
# Plot histograms of each features distribution
hist = non_poi_df.hist(column=['exercised_stock_options', 'bonus', 'restricted_stock', 'shared_receipt_with_poi'],
                        figsize=(18,3), layout=(1,4))
```




```
In [21]: # POI
# Plot histograms of each features distribution
hist = poi_df.hist(column=['exercised_stock_options', 'bonus', 'restricted_stock', 'shared_receipt_with_poi'],
                    figsize=(18,3), layout=(1,4))
```




```
In [22]: # Non POI
# Plot histograms of each features distribution
hist = non_poi_df.hist(column=['restricted_stock_deferred', 'total_stock_value', 'expenses', 'loan_advances'],
                        figsize=(18,3), layout=(1,4))
```




```
In [23]: # POI
# Plot histograms of each features distribution
hist = poi_df.hist(column=['restricted_stock_deferred', 'total_stock_value',
'expenses', 'loan_advances'],
figsize=(18,3), layout=(1,4))
```




```
In [24]: # Non POI
# Plot histograms of each features distribution
hist = non_poi_df.hist(column=['from_messages', 'other', 'from_this_person_to_poi',
'poi'],
figsize=(18,3), layout=(1,4))
```




```
In [25]: # POI
# Plot histograms of each features distribution
hist = poi_df.hist(column=['from_messages', 'other', 'from_this_person_to_poi',
'poi'],
figsize=(18,3), layout=(1,4))
```



```
In [26]: # Non POI
# Plot histograms of each features distribution
hist = non_poi_df.hist(column=['director_fees', 'deferred_income', 'long_term_incentive', 'from_poi_to_this_person'],
                        figsize=(18,3), layout=(1,4))
```




```
In [27]: # POI
# Plot histograms of each features distribution
hist = poi_df.hist(column=['director_fees', 'deferred_income', 'long_term_incentive', 'from_poi_to_this_person'],
                    figsize=(18,3), layout=(1,4))
```




Feature Selection, Scaling, Engineering

```
In [28]: # Make new columns in df for new features email features
df['%_from_poi'] = (df['from_poi_to_this_person'] / df['to_messages'])*100
df['%_to_poi'] = (df['from_this_person_to_poi'] / df['from_messages'])*100
df['%_shared_receipt_with_poi'] = (df['shared_receipt_with_poi'] / df['to_messages'])*100
```

```
In [29]: # Plot histograms of each new features distribution
hist = df.hist(column=['%_from_poi', '%_to_poi', '%_shared_receipt_with_poi'],
                figsize=(18,5), layout=(1,3))
```



```
In [30]: # Update poi_df and plot histograms of each new features distribution for pois
poi_df = df[df['poi'] == True]
hist = poi_df.hist(column=['%_from_poi', '%_to_poi', '%_shared_receipt_with_poi'], figsize=(18,5), layout=(1,3))
```



Preprocessing

Save Data to New Dict, Extract Features and Labels, Replace NaNs, Select Features, Split Training and Testing Sets, Balance Classes for Training

```
In [31]: # Replace nans with medians
df = df.fillna(df.median())
```

Missing values must be dealt with prior to building a classifier so imputing is used to replace any missing values with the median value for that feature. The median is used in this case rather than mean as it is less susceptible to outliers and there are known outliers in this dataset.

```
In [32]: # Convert df to dict and preview
my_dataset = df.to_dict('index')
```

```
In [33]: # Select what features to use. Features_list is a list of strings, each of which is a feature name.
# The first feature must be "poi".
features_list = ['poi', 'salary', 'total_payments', 'exercised_stock_options', 'bonus', 'restricted_stock', 'total_stock_value', 'expenses', 'other', 'deferred_income', 'long_term_incentive', '%_from_poi', '%_to_poi', '%_shared_receipt_with_poi']
```

Features with absolute values (to_mesages, from_messages, from_this_person_to_poi, from_poi_to_this_person, and shared_receipt_with_poi) were dropped in favor of relative values (%_from_poi, %_to_poi, %_shared_receipt_with_poi). Features with no non-null values for pois (director_fees and restricted_stock_deferred) and features with more null than non_null values for pois (deferral_payments and loan_advances) were also dropped. The only string type feature (email_address) was dropped as well.

```
In [34]: # Extract features and labels from dataset for local testing
data = featureFormat(my_dataset, features_list, sort_keys = True)
labels, features = targetFeatureSplit(data)
```

```
In [35]: # Scale with MinMaxScaler
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaler.fit(features)
features = scaler.transform(features)
```

The feature values are scaled to a range of 0-1 with MinMaxScaler to deal with the variation in magnitude, sign, and units of the data. This is not necessary for many of the algorithms used, but is essential for several (those like KNN that use Euclidean distance) and should not have any negative effect for the other algorithms.

```
In [36]: # Use select k best to narrow chosen features down to best 10
from sklearn.feature_selection import SelectKBest
selector = SelectKBest(k=10).fit(features, labels)
feat_nums = selector.get_support(indices=True)
features = selector.transform(features)

# Print scores for all features, print scores for 10 chosen features, and print names of 10 chosen features
features_array = np.array(features_list[1:])
print("All Scores:", selector.scores_)
print("Top 10 Scores:", selector.scores_[feat_nums])
print("Top 10 Features:", features_array[feat_nums])

# Make new features list
kfeatures_list = list(features_array[feat_nums])
features_list = ['poi'] + kfeatures_list

('All Scores:', array([11.06405906,  8.50745309, 27.67840142, 16.13583336,
 8.58242675,
      23.92024067,  1.03887138,  3.97611644, 10.47995347,  8.58240739,
 1.91873384, 16.19447573,  6.6525593 ]))
('Top 10 Scores:', array([11.06405906,  8.50745309, 27.67840142, 16.13583336,
 8.58242675,
      23.92024067, 10.47995347,  8.58240739, 16.19447573,  6.6525593 ]))
('Top 10 Features:', array(['salary', 'total_payments', 'exercised_stock_opti
ons', 'bonus',
      'restricted_stock', 'total_stock_value', 'deferred_income',
      'long_term_incentive', '%_to_poi', '%_shared_receipt_with_poi'],
      dtype='<S25'))
```

The feature set was narrowed down manually to 13 features but it can be hard to predict usefulness of a feature without the help of an algorithm. SelectKBest is used to determine the predictive potential of each remaining feature. Since there are 3 features with very low scores (less than 5) these 3 are dropped and the 10 most useful are retained.


```
In [37]: # Split training and testing data for validation
from sklearn.model_selection import train_test_split
features_train, features_test, labels_train, labels_test = \
    train_test_split(features, labels, test_size=0.3, random_state=42)
```

```
In [38]: # Make data frames of training labels and training features, concat to single
         # training df, set column labels to features,
         # and preview result
train_ls = pd.DataFrame(labels_train)
train_fs = pd.DataFrame(features_train)
train = pd.concat([train_ls, train_fs], axis=1)
train.columns = features_list
```

```
In [39]: # Make data frames of testing labels and testing features, concat to single te
         # sting df, set column labels to features,
         # and preview result
test_ls = pd.DataFrame(labels_test)
test_fs = pd.DataFrame(features_test)
test = pd.concat([test_ls, test_fs], axis=1)
test.columns = features_list
```

```
In [40]: # Split training df into non-pois (majority) and pois (minority), then use ran
         # dom upsampling on pois to balance classes,
         # then recombine into one training df
df_maj = train[train.poi==0]
df_min = train[train.poi==1]

from sklearn.utils import resample

df_min_up = resample(df_min, replace = True, n_samples = len(df_maj), random_s
tate = 0)

train = pd.concat([df_maj, df_min_up])
train = train.reset_index(drop = True)
```

Since there is a pretty heavy class imbalance in the data with only 18 POIs out of 145 total observations (about a 5:95 split) random upsampling is used to artificially inflate the training set with more POI observations. This balancing of classes in the training set should give better predictive power in the trained models.

```
In [41]: # Make full df that contains both the upsampled training df and the testing df
full = pd.concat([train, test])
full = full.reset_index(drop = True)
```

```
In [42]: # Convert training df to dict
training_set = train.to_dict('index')
```

```
In [43]: # Extract features and labels from training dataset
data = featureFormat(training_set, features_list, sort_keys = True)
labels_train, features_train = targetFeatureSplit(data)
```

```
In [44]: # Convert full df to dict and preview
my_dataset = full.to_dict('index')
```

```
In [45]: # Extract features and labels from full dataset
data = featureFormat(my_dataset, features_list, sort_keys = True)
labels, features = targetFeatureSplit(data)
```

After preprocessing, the data contains no missing values, only has the 10 most useful features, and is scaled to a range of 0-1. The data is organized as: labels_train and features_train (a training set of a features array and a labels array with classes balanced by random upsampling of the minority class), labels_test and features_test (a testing set of a features array and a labels array set aside for validation containing 30% of the original data), and features and labels (a features array and a labels array containing the combination of the training and testing sets).

Algorithm Selection

Begin by making classifiers with default hyperparameters. Some of the more common classification algorithms will be tested including: Gaussian Naive Bayes, Random Forest, AdaBoost, Gradient Boost, Decision Tree, Logistic Regression, K Neighbors, and Support Vector Machine.

```
In [76]: # Make a classifier with default parameters and predict to see initial results
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score

clf = GaussianNB()
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
GNBclf = clf

('Accuracy Score:', 0.8863636363636364)
('Precision Score:', 0.4)
('Recall Score:', 0.5)
('F1 Score:', 0.4444444444444445)
```

```
In [77]: # Make a classifier with default parameters and predict to see initial results
from sklearn.ensemble import RandomForestClassifier

clf = RandomForestClassifier(random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
RFclf = clf

('Accuracy Score:', 0.8409090909090909)
('Precision Score:', 0.2857142857142857)
('Recall Score:', 0.5)
('F1 Score:', 0.36363636363636365)
```

```
In [78]: # Make a classifier with default parameters and predict to see initial results
from sklearn.ensemble import AdaBoostClassifier

clf = AdaBoostClassifier(random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
ABclf = clf

('Accuracy Score:', 0.8181818181818182)
('Precision Score:', 0.25)
('Recall Score:', 0.5)
('F1 Score:', 0.3333333333333333)
```

```
In [79]: # Make a classifier with default parameters and predict to see initial results
from sklearn.ensemble import GradientBoostingClassifier

clf = GradientBoostingClassifier(random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
GBclf = clf

('Accuracy Score:', 0.9090909090909091)
('Precision Score:', 0.5)
('Recall Score:', 0.5)
('F1 Score:', 0.5)
```

```
In [80]: # Make a classifier with default parameters and predict to see initial results
from sklearn.tree import DecisionTreeClassifier

clf = DecisionTreeClassifier(random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
DTclf = clf

('Accuracy Score:', 0.8863636363636364)
('Precision Score:', 0.4)
('Recall Score:', 0.5)
('F1 Score:', 0.4444444444444445)
```

```
In [81]: # Make a classifier with default parameters and predict to see initial results
from sklearn.linear_model import LogisticRegression

clf = LogisticRegression(random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
LRclf = clf

('Accuracy Score:', 0.7954545454545454)
('Precision Score:', 0.2727272727272727)
('Recall Score:', 0.75)
('F1 Score:', 0.39999999999999997)
```

```
In [82]: # Make a classifier with default parameters and predict to see initial results
from sklearn.neighbors import KNeighborsClassifier

clf = KNeighborsClassifier()
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
KNclf = clf

('Accuracy Score:', 0.7954545454545454)
('Precision Score:', 0.2222222222222222)
('Recall Score:', 0.5)
('F1 Score:', 0.30769230769230765)
```

```
In [83]: # Make a classifier with default parameters and predict to see initial results
from sklearn.svm import SVC

clf = SVC(random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
SVMclf = clf

('Accuracy Score:', 0.8863636363636364)
('Precision Score:', 0.3333333333333333)
('Recall Score:', 0.25)
('F1 Score:', 0.28571428571428575)
```

```
In [84]: # Make Lists of clfs and attributes then make into df
clfs = [GNBclf, RFclf, ABclf, GBclf, DTclf, LRclf, KNCclf, SVMclf]
clf_names = ['GNBclf', 'RFclf', 'ABclf', 'GBclf', 'DTclf', 'LRclf', 'KNCclf',
'SVMclf']
clf_details = [str(i) for i in clfs]
clf_acc = []
clf_pre = []
clf_rec = []
clf_f1 = []

for clf in clfs:
    clf_acc.append(accuracy_score(labels_test, clf.predict(features_test)))
    clf_pre.append(precision_score(labels_test, clf.predict(features_test)))
    clf_rec.append(recall_score(labels_test, clf.predict(features_test)))
    clf_f1.append(f1_score(labels_test, clf.predict(features_test)))

clf_df = pd.DataFrame(list(zip(clf_acc, clf_pre, clf_rec, clf_f1, clf_details
)),
                        index = clf_names, columns = ['accuracy', 'precision',
'recall', 'f1', 'details'])

clf_df
```

Out[84]:

	accuracy	precision	recall	f1	details
GNBclf	0.886364	0.400000	0.50	0.444444	GaussianNB(priors=None, var_smoothing=1e-09)
RFclf	0.840909	0.285714	0.50	0.363636	RandomForestClassifier(bootstrap=True, class_w...
ABclf	0.818182	0.250000	0.50	0.333333	AdaBoostClassifier(algorithm='SAMME.R', base_e...
GBclf	0.909091	0.500000	0.50	0.500000	GradientBoostingClassifier(criterion='friedman...
DTclf	0.886364	0.400000	0.50	0.444444	DecisionTreeClassifier(class_weight=None, crit...
LRclf	0.795455	0.272727	0.75	0.400000	LogisticRegression(C=1.0, class_weight=None, d...
KNCclf	0.795455	0.222222	0.50	0.307692	KNeighborsClassifier(algorithm='auto', leaf_si...
SVMclf	0.886364	0.333333	0.25	0.285714	SVC(C=1.0, cache_size=200, class_weight=None, ...

Several classifiers give pretty good results with the default hyperparameters. Gridsearchcv will be used on all to determine the optimal hyperparameters for each (except GaussianNB since there is not really anything to tune in that case).

```
In [104]: # Tune Random Forest to optimize F1 and make predictions with best clf
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import PredefinedSplit
params = {'n_estimators' : (3, 4, 5, 10, 20), 'min_samples_split' : (2, 3, 4,
5, 6), 'random_state' : np.arange(0,100)}
alg = RandomForestClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'f1', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
RFclf_f = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
```

```
('Best Parameters:', {'min_samples_split': 2, 'n_estimators': 4, 'random_state': 91})
```

```
('Accuracy Score:', 0.9318181818181818)
```

```
('Precision Score:', 1.0)
```

```
('Recall Score:', 0.25)
```

```
('F1 Score:', 0.4)
```

```
C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
```

```
DeprecationWarning)
```

```
In [105]: # Tune Random Forest to optimize Recall and make predictions with best clf
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import PredefinedSplit
params = {'n_estimators' : (3, 4, 5, 10, 20), 'min_samples_split' : (2, 3, 4, 5, 6), 'random_state' : np.arange(0,100)}
alg = RandomForestClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'recall', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
RFclf_r = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))

('Best Parameters:', {'min_samples_split': 2, 'n_estimators': 3, 'random_state': 2})
('Accuracy Score:', 0.8636363636363636)
('Precision Score:', 0.3333333333333333)
('Recall Score:', 0.5)
('F1 Score:', 0.4)

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
  DeprecationWarning)
```

```
In [106]: # Tune AdaBoost to optimize F1 and make predictions with best clf
params = {'n_estimators' : (10, 50, 100, 500, 1000), 'learning_rate' : (0.001, 0.01, 0.1, 0.5, 1),
          'random_state' : np.arange(0,100)}
alg = AdaBoostClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'f1', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
ABclf_f = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
  DeprecationWarning)

('Best Parameters:', {'n_estimators': 100, 'learning_rate': 0.5, 'random_state': 0})
('Accuracy Score:', 0.8863636363636364)
('Precision Score:', 0.4)
('Recall Score:', 0.5)
('F1 Score:', 0.4444444444444445)
```

```
In [107]: # Tune AdaBoost to optimize Recall and make predictions with best clf
params = {'n_estimators' : (10, 50, 100, 500, 1000), 'learning_rate' : (0.001,
0.01, 0.1, 0.5, 1),
          'random_state' : np.arange(0,100)}
alg = AdaBoostClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'recall', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
ABclf_r = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))

('Best Parameters:', {'n_estimators': 100, 'learning_rate': 0.1, 'random_state': 0})
('Accuracy Score:', 0.7954545454545454)
('Precision Score:', 0.2222222222222222)
('Recall Score:', 0.5)
('F1 Score:', 0.30769230769230765)

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)
```

```
In [108]: # Tune Gradient Boost to optimize F1 and make predictions with best clf
params = {'n_estimators' : (10, 50, 100, 200, 500), 'learning_rate' : (0.0001,
0.001, 0.01, 0.1, 0.5, 1),
          'random_state' : np.arange(0,100)}
alg = GradientBoostingClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'f1', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
GBclf_f = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))

('Best Parameters:', {'n_estimators': 50, 'learning_rate': 1, 'random_state': 12})
('Accuracy Score:', 0.9090909090909091)
('Precision Score:', 0.5)
('Recall Score:', 0.5)
('F1 Score:', 0.5)

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)
```



```
In [109]: # Tune Gradient Boost to optimize Recall and make predictions with best clf
params = {'n_estimators' : (10, 50, 100, 200, 500), 'learning_rate' : (0.0001,
0.001, 0.01, 0.1, 0.5, 1),
          'random_state' : np.arange(0,100)}
alg = GradientBoostingClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'recall', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
GBclf_r = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
```

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

```
('Best Parameters:', {'n_estimators': 500, 'learning_rate': 0.01, 'random_state': 0})
('Accuracy Score:', 0.8636363636363636)
('Precision Score:', 0.3333333333333333)
('Recall Score:', 0.5)
('F1 Score:', 0.4)
```

```
In [110]: # Tune Decision Tree to optimize F1 and make predictions with best clf
params = {'min_samples_split' : (np.arange(2,11)), 'random_state' : np.arange(
0,100)}
alg = DecisionTreeClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'f1', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
DTclf_f = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
```

```
('Best Parameters:', {'min_samples_split': 2, 'random_state': 88})
('Accuracy Score:', 0.9090909090909091)
('Precision Score:', 0.5)
('Recall Score:', 0.5)
('F1 Score:', 0.5)
```

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

```
In [111]: # Tune Decision Tree to optimize Recall and make predictions with best clf
params = {'min_samples_split' : (np.arange(2,11)), 'random_state' : np.arange(
0,100)}
alg = DecisionTreeClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'recall', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
DTclf_r = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
```

```
('Best Parameters:', {'min_samples_split': 2, 'random_state': 0})
('Accuracy Score:', 0.8863636363636364)
('Precision Score:', 0.4)
('Recall Score:', 0.5)
('F1 Score:', 0.44444444444444445)
```

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

```
In [112]: # Tune Logistic Regression to optimize F1 and make predictions with best clf
params = {'penalty' : ('l1', 'l2'), 'C' : (np.logspace(-4, 4, 50)), 'random_state' : np.arange(0,100)}
alg = LogisticRegression()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'f1', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
LRclf_f = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
```

```
('Best Parameters:', {'penalty': 'l2', 'C': 11.513953993264456, 'random_state': 0})
('Accuracy Score:', 0.7727272727272727)
('Precision Score:', 0.2)
('Recall Score:', 0.5)
('F1 Score:', 0.28571428571428575)
```

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

```
In [113]: # Tune Logistic Regression to optimize Recall and make predictions with best c
          lf
          params = {'penalty' : ('l1', 'l2'), 'C' : (np.logspace(-4, 4, 50)), 'random_st
          ate' : np.arange(0,100)}
          alg = LogisticRegression()
          clf = GridSearchCV(alg, params, cv = 5, scoring = 'recall', n_jobs = -1)
          clf.fit(features_train, labels_train)
          print("Best Parameters:", clf.best_params_)
          pred = clf.predict(features_test)
          LRclf_r = clf
          print("Accuracy Score:", accuracy_score(labels_test, pred))
          print("Precision Score:", precision_score(labels_test, pred))
          print("Recall Score:", recall_score(labels_test, pred))
          print("F1 Score:", f1_score(labels_test, pred))
```

```
('Best Parameters:', {'penalty': 'l2', 'C': 0.0001, 'random_state': 0})
('Accuracy Score:', 0.1590909090909091)
('Precision Score:', 0.0975609756097561)
('Recall Score:', 1.0)
('F1 Score:', 0.17777777777777776)
```

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

```
In [114]: # Tune KNN to optimize F1 and make predictions with best clf
          params = {'n_neighbors' : (10, 20, 30, 40), 'p' : (1, 2), 'leaf_size' : (2, 3,
          4, 5, 10, 20)}
          alg = KNeighborsClassifier()
          clf = GridSearchCV(alg, params, cv = 5, scoring = 'f1', n_jobs = -1)
          clf.fit(features_train, labels_train)
          print("Best Parameters:", clf.best_params_)
          pred = clf.predict(features_test)
          KNclf_f = clf
          print("Accuracy Score:", accuracy_score(labels_test, pred))
          print("Precision Score:", precision_score(labels_test, pred))
          print("Recall Score:", recall_score(labels_test, pred))
          print("F1 Score:", f1_score(labels_test, pred))
```

```
('Best Parameters:', {'n_neighbors': 10, 'leaf_size': 2, 'p': 2})
('Accuracy Score:', 0.7954545454545454)
('Precision Score:', 0.2222222222222222)
('Recall Score:', 0.5)
('F1 Score:', 0.30769230769230765)
```

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

```
In [115]: # Tune KNN to optimize Recall and make predictions with best clf
params = {'n_neighbors' : (10, 20, 30, 40), 'p' : (1, 2), 'leaf_size' : (2, 3, 4, 5, 10, 20)}
alg = KNeighborsClassifier()
clf = GridSearchCV(alg, params, cv = 5, scoring = 'recall', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
KNclf_r = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
```

```
('Best Parameters:', {'n_neighbors': 10, 'leaf_size': 2, 'p': 2})
('Accuracy Score:', 0.7954545454545454)
('Precision Score:', 0.2222222222222222)
('Recall Score:', 0.5)
('F1 Score:', 0.30769230769230765)
```

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

```
In [116]: # Tune SVM to optimize F1 and make predictions with best clf
params = {'C' : (0.1, 1, 10, 100), 'kernel' : ('rbf', 'linear', 'poly', 'sigmoid'),
          'gamma' : ('auto', 0.01, 1, 10)}
alg = SVC(random_state = 1)
clf = GridSearchCV(alg, params, cv = 5, scoring = 'f1', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
SVMclf_f = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))
```

```
('Best Parameters:', {'kernel': 'rbf', 'C': 10, 'gamma': 10})
('Accuracy Score:', 0.8409090909090909)
('Precision Score:', 0.2857142857142857)
('Recall Score:', 0.5)
('F1 Score:', 0.36363636363636365)
```

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
DeprecationWarning)

```

In [117]: # Tune SVM to optimize Recall and make predictions with best clf
params = {'C' : (0.1, 1, 10, 100), 'kernel' : ('rbf', 'linear', 'poly', 'sigmoid'),
          'gamma' : ('auto', 0.01, 1, 10)}
alg = SVC(random_state = 1)
clf = GridSearchCV(alg, params, cv = 5, scoring = 'recall', n_jobs = -1)
clf.fit(features_train, labels_train)
print("Best Parameters:", clf.best_params_)
pred = clf.predict(features_test)
SVMclf_r = clf
print("Accuracy Score:", accuracy_score(labels_test, pred))
print("Precision Score:", precision_score(labels_test, pred))
print("Recall Score:", recall_score(labels_test, pred))
print("F1 Score:", f1_score(labels_test, pred))

('Best Parameters:', {'kernel': 'poly', 'C': 1, 'gamma': 10})
('Accuracy Score:', 0.7954545454545454)
('Precision Score:', 0.2222222222222222)
('Recall Score:', 0.5)
('F1 Score:', 0.30769230769230765)

C:\Users\rbruton\.conda\envs\Python 2 2\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.
  DeprecationWarning)

```

```
In [118]: # Make Lists of clfs and attributes then make into df for F1 optimized clfs
clfs = [GNBclf, RFclf_f, RFclf_r, ABclf_f, ABclf_r, GBclf_f, GBclf_r, DTclf_f,
        DTclf_r, LRclf_f, LRclf_r,
        KNclf_f, KNclf_r, SVMclf_f, SVMclf_r]
clf_names = ['GNBclf', 'RFclf_f', 'RFclf_r', 'ABclf_f', 'ABclf_r', 'GBclf_f',
             'GBclf_r', 'DTclf_f', 'DTclf_r',
             'LRclf_f', 'LRclf_r', 'KNclf_f', 'KNclf_r', 'SVMclf_f', 'SVMclf_r']
clf_details = [str(i) for i in clfs]
clf_acc = []
clf_pre = []
clf_rec = []
clf_f1 = []

for clf in clfs:
    clf_acc.append(accuracy_score(labels_test, clf.predict(features_test)))
    clf_pre.append(precision_score(labels_test, clf.predict(features_test)))
    clf_rec.append(recall_score(labels_test, clf.predict(features_test)))
    clf_f1.append(f1_score(labels_test, clf.predict(features_test)))

clf_df = pd.DataFrame(list(zip(clf_acc, clf_pre, clf_rec, clf_f1, clf_details)),
                      index = clf_names, columns = ['accuracy', 'precision',
                                                    'recall', 'f1', 'details'])

clf_df
```

Out[118]:

	accuracy	precision	recall	f1	details
GNBclf	0.886364	0.400000	0.50	0.444444	GaussianNB(priors=None, var_smoothing=1e-09)
RFclf_f	0.931818	1.000000	0.25	0.400000	GridSearchCV(cv=5, error_score='raise-deprecat...
RFclf_r	0.863636	0.333333	0.50	0.400000	GridSearchCV(cv=5, error_score='raise-deprecat...
ABclf_f	0.886364	0.400000	0.50	0.444444	GridSearchCV(cv=5, error_score='raise-deprecat...
ABclf_r	0.795455	0.222222	0.50	0.307692	GridSearchCV(cv=5, error_score='raise-deprecat...
GBclf_f	0.909091	0.500000	0.50	0.500000	GridSearchCV(cv=5, error_score='raise-deprecat...
GBclf_r	0.863636	0.333333	0.50	0.400000	GridSearchCV(cv=5, error_score='raise-deprecat...
DTclf_f	0.909091	0.500000	0.50	0.500000	GridSearchCV(cv=5, error_score='raise-deprecat...
DTclf_r	0.886364	0.400000	0.50	0.444444	GridSearchCV(cv=5, error_score='raise-deprecat...
LRclf_f	0.772727	0.200000	0.50	0.285714	GridSearchCV(cv=5, error_score='raise-deprecat...
LRclf_r	0.159091	0.097561	1.00	0.177778	GridSearchCV(cv=5, error_score='raise-deprecat...
KNclf_f	0.795455	0.222222	0.50	0.307692	GridSearchCV(cv=5, error_score='raise-deprecat...
KNclf_r	0.795455	0.222222	0.50	0.307692	GridSearchCV(cv=5, error_score='raise-deprecat...
SVMclf_f	0.840909	0.285714	0.50	0.363636	GridSearchCV(cv=5, error_score='raise-deprecat...
SVMclf_r	0.795455	0.222222	0.50	0.307692	GridSearchCV(cv=5, error_score='raise-deprecat...

Many classifiers give good results during the gridsearchcv optimization so all classifiers with optimized hyperparameters will be tested under normal conditions.

```
In [119]: # Optimized for F1
# Make a Random Forest classifier with optimized parameters and test
clf = RandomForestClassifier(random_state = 91, min_samples_split = 2, n_estimators = 4)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
RFclf_f = clf

# Make an AdaBoost classifier with optimized parameters and test
clf = AdaBoostClassifier(n_estimators = 100, learning_rate = 0.5, random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
ABclf_f = clf

# Make a Gradient Boost classifier with optimized parameters and test
clf = GradientBoostingClassifier(random_state = 12, learning_rate = 1, n_estimators = 50)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
GBclf_f = clf

# Make a Decision Tree classifier with optimized parameters and test
clf = DecisionTreeClassifier(random_state = 88, min_samples_split = 2)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
DTclf_f = clf

# Make a Logistic Regression classifier with optimized parameters and test
clf = LogisticRegression(penalty = 'l2', C = 11.5, random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
LRclf_f = clf

# Make a K Neighbors classifier with optimized parameters and test
clf = KNeighborsClassifier(n_neighbors = 10, leaf_size = 2, p = 2)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
KNclf_f = clf

# Make a SVM classifier with optimized parameters and test
clf = SVC(kernel = 'rbf', C = 10, gamma = 10)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
SVMclf_f = clf
```

```
In [121]: # Optimized for Recall
# Make a Random Forest classifier with optimized parameters and test
clf = RandomForestClassifier(random_state = 2, min_samples_split = 2, n_estimators = 3)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
RFclf_r = clf

# Make an AdaBoost classifier with optimized parameters and test
clf = AdaBoostClassifier(n_estimators = 100, learning_rate = 0.1, random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
ABclf_r = clf

# Make a Gradient Boost classifier with optimized parameters and test
clf = GradientBoostingClassifier(random_state = 0, learning_rate = 0.01, n_estimators = 500)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
GBclf_r = clf

# Make a Decision Tree classifier with optimized parameters and test
clf = DecisionTreeClassifier(random_state = 0, min_samples_split = 2)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
DTclf_r = clf

# Make a Logistic Regression classifier with optimized parameters and test
clf = LogisticRegression(penalty = 'l2', C = 0.0001, random_state = 0)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
LRclf_r = clf

# Make a K Neighbors classifier with optimized parameters and test
clf = KNeighborsClassifier(n_neighbors = 10, leaf_size = 2, p = 2)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
KNclf_r = clf

# Make a SVM classifier with optimized parameters and test
clf = SVC(kernel = 'poly', C = 1, gamma = 10)
clf.fit(features_train, labels_train)
pred = clf.predict(features_test)
SVMclf_r = clf
```



```

In [122]: # Make Lists of clfs and attributes then make into df for F1 and Recall optimi
zed clfs
clfs = [GNBclf, RFclf_f, RFclf_r, ABclf_f, ABclf_r, GBclf_f, GBclf_r, DTclf_f,
DTclf_r, LRclf_f, LRclf_r,
        KNclf_f, KNclf_r, SVMclf_f, SVMclf_r]
clf_names = ['GNBclf', 'RFclf_f', 'RFclf_r', 'ABclf_f', 'ABclf_r', 'GBclf_f',
'GBclf_r', 'DTclf_f', 'DTclf_r',
            'LRclf_f', 'LRclf_r', 'KNclf_f', 'KNclf_r', 'SVMclf_f', 'SVMclf_
r']
clf_details = [str(i) for i in clfs]
clf_acc = []
clf_pre = []
clf_rec = []
clf_f1 = []

for clf in clfs:
    clf_acc.append(accuracy_score(labels_test, clf.predict(features_test)))
    clf_pre.append(precision_score(labels_test, clf.predict(features_test)))
    clf_rec.append(recall_score(labels_test, clf.predict(features_test)))
    clf_f1.append(f1_score(labels_test, clf.predict(features_test)))

clf_df = pd.DataFrame(list(zip(clf_acc, clf_pre, clf_rec, clf_f1, clf_details
)),
                      index = clf_names, columns = ['accuracy', 'precision',
'recall', 'f1', 'details'])

clf_df

```

Out[122]:

	accuracy	precision	recall	f1	details
GNBclf	0.886364	0.400000	0.50	0.444444	GaussianNB(priors=None, var_smoothing=1e-09)
RFclf_f	0.931818	1.000000	0.25	0.400000	RandomForestClassifier(bootstrap=True, class_w...
RFclf_r	0.863636	0.333333	0.50	0.400000	RandomForestClassifier(bootstrap=True, class_w...
ABclf_f	0.886364	0.400000	0.50	0.444444	AdaBoostClassifier(algorithm='SAMME.R', base_e...
ABclf_r	0.795455	0.222222	0.50	0.307692	AdaBoostClassifier(algorithm='SAMME.R', base_e...
GBclf_f	0.909091	0.500000	0.50	0.500000	GradientBoostingClassifier(criterion='friedman...
GBclf_r	0.863636	0.333333	0.50	0.400000	GradientBoostingClassifier(criterion='friedman...
DTclf_f	0.909091	0.500000	0.50	0.500000	DecisionTreeClassifier(class_weight=None, crit...
DTclf_r	0.886364	0.400000	0.50	0.444444	DecisionTreeClassifier(class_weight=None, crit...
LRclf_f	0.772727	0.200000	0.50	0.285714	LogisticRegression(C=11.5, class_weight=None, ...
LRclf_r	0.159091	0.097561	1.00	0.177778	LogisticRegression(C=0.0001, class_weight=None...
KNclf_f	0.795455	0.222222	0.50	0.307692	KNeighborsClassifier(algorithm='auto', leaf_si...
KNclf_r	0.795455	0.222222	0.50	0.307692	KNeighborsClassifier(algorithm='auto', leaf_si...
SVMclf_f	0.840909	0.285714	0.50	0.363636	SVC(C=10, cache_size=200, class_weight=None, c...
SVMclf_r	0.795455	0.222222	0.50	0.307692	SVC(C=1, cache_size=200, class_weight=None, co...

The Gradient Boosting classifier with `random_state = 12`, `learning_rate = 1`, and `n_estimators = 50` gives the best results and will be the final choice of classifier.

```
In [127]: # Assign best performing classifier to clf
         clf = GBclf_f
```

```
In [128]: ### Task 6: Dump your classifier, dataset, and features_list so anyone can
         ### check your results. You do not need to change anything below, but make sur
         e
         ### that the version of poi_id.py that you submit can be run on its own and
         ### generates the necessary .pkl files for validating your results.

         dump_classifier_and_data(clf, my_dataset, features_list)
```

Output from Tester.py

GradientBoostingClassifier(criterion='friedman_mse', init=None, learning_rate=1, loss='deviance', max_depth=3, max_features=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=50, n_iter_no_change=None, presort='auto', random_state=12, subsample=1.0, tol=0.0001, validation_fraction=0.1, verbose=0, warm_start=False)

Accuracy: 0.93218 Precision: 0.88065 Recall: 0.96500 F1: 0.92090 F2: 0.94686

Total predictions: 22000

True positives: 8685 False positives: 1177 False negatives: 315 True negatives: 11823

```
In [ ]:
```